


Fundamentals of Database Systems

Seventh Edition



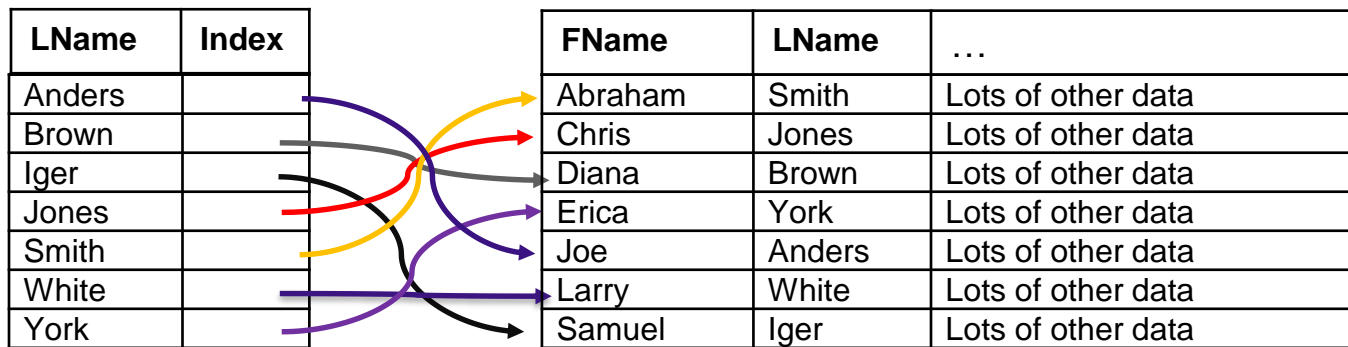
Chapter 17

Indexing Structures for
Files and Physical
Database Design

Introduction

- Indexes used to speed up record retrieval in response to certain search conditions
- Index structures provide secondary access paths
- Any field can be used to create an index
 - Multiple indexes can be constructed
- Most indexes based on ordered files
 - Tree data structures organize the index

Index example



17.1 Types of Single-Level Ordered Indexes

- Ordered index similar to index in a textbook
- Indexing field (attribute)
 - Index stores each value of the index field with list of pointers to all disk blocks that contain records with that field value
- Values in index are ordered
- Primary index
 - Specified on the ordering key field of ordered file of records

17.1 Types of Single-Level Ordered Indexes

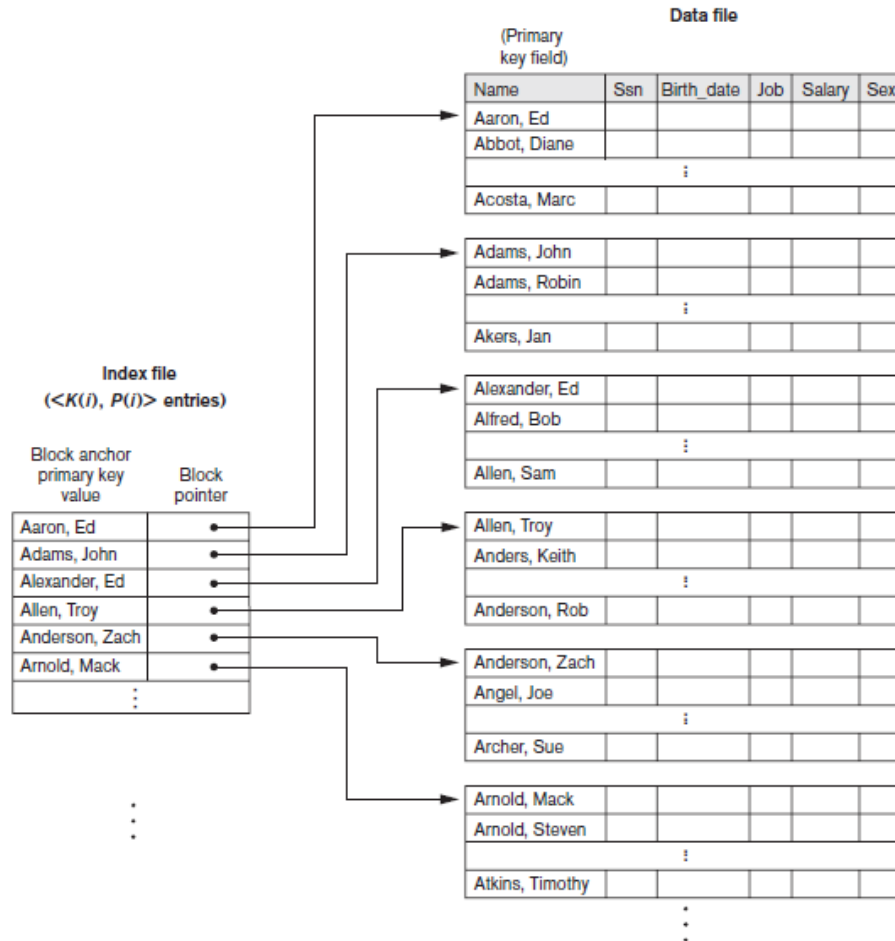
- Clustering index
 - Used if numerous records can have the same value for the ordering field
- Secondary index
 - Can be specified on any nonordering field
 - Data file can have several secondary indexes

Primary Indexes

- Ordered file with two fields
 - Primary key, $K(i)$
 - Pointer to a disk block, $P(i)$
- One index entry in the index file for each block in the data file
- Indexes may be dense or sparse
 - Dense index has an index entry for every search key value in the data file
 - Sparse index has entries for only some search values

Primary Indexes

Figure 17.1 Primary index on the ordering key field of the file shown in Figure 16.7



Primary Indexes

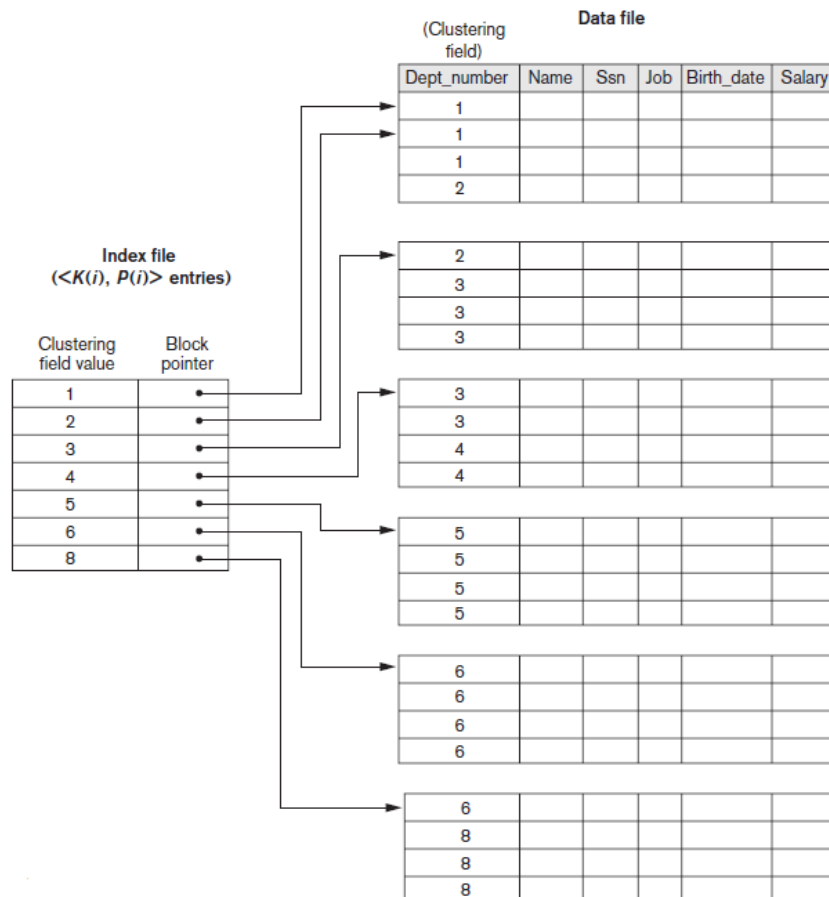
- Major problem: insertion and deletion of records
 - Move records around and change index values
 - Solutions
 - Use unordered overflow file
 - Use linked list of overflow records

Clustering Indexes

- Clustering field
 - File records are physically ordered on a nonkey field without a distinct value for each record
- Ordered file with two fields
 - Same type as clustering field
 - Disk block pointer

Clustering Indexes

Figure 17.2 A clustering index on the Dept_number ordering nonkey field of an EMPLOYEE file

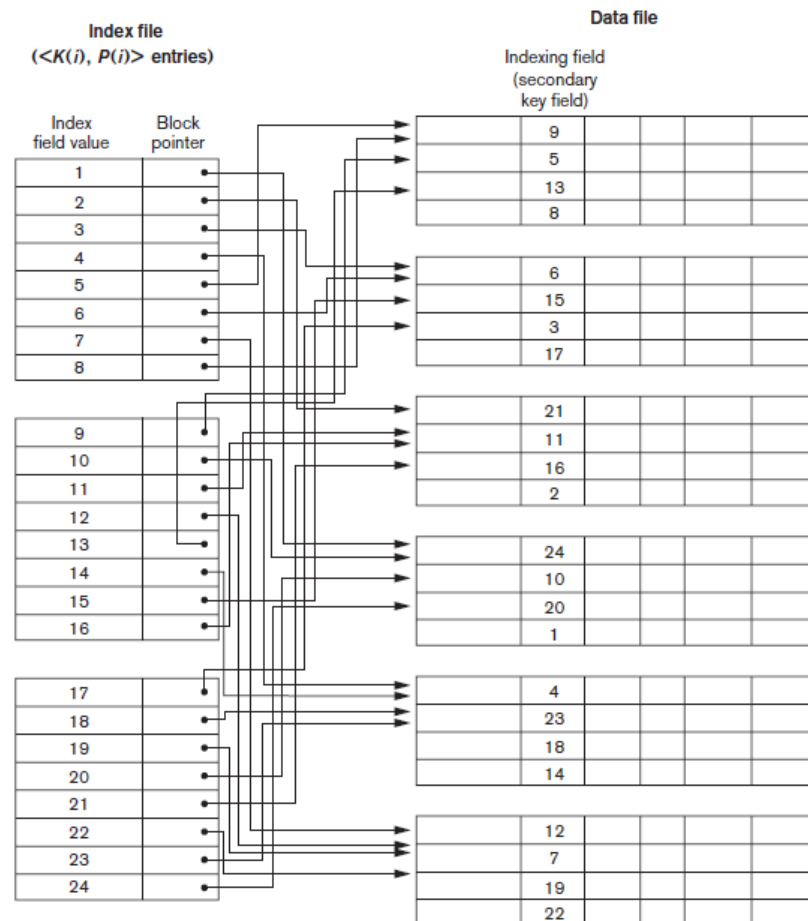


Secondary Indexes

- Provide secondary means of accessing a data file
 - Some primary access exists
- Ordered file with two fields
 - Indexing field, $K(i)$
 - Block pointer or record pointer, $P(i)$
- Usually need more storage space and longer search time than primary index
 - Improved search time for arbitrary record

Secondary Indexes

Figure 17.4 Dense secondary index (with block pointers) on a nonordering key field of a file.



17.1 Types of Single-Level Ordered Indexes

Table 17.1 Types of indexes based on the properties of the indexing field

| | Index Field Used for Physical Ordering of the File | Index Field Not Used for Physical Ordering of the File |
|--------------------------|---|---|
| Indexing field is key | Primary index | Secondary index (Key) |
| Indexing field is nonkey | Clustering index | Secondary index (NonKey) |

17.1 Types of Single-Level Ordered Indexes

Table 17.2 Properties of Index Types

| Type of Index | Number of (First-Level) Index Entries | Dense or Nondense (Sparse) | Block Anchoring on the Data File |
|--------------------|--|----------------------------|----------------------------------|
| Primary | Number of blocks in data file | Nondense | Yes |
| Clustering | Number of distinct index field values | Nondense | Yes / no ^a |
| Secondary (key) | Number of records in data file | Dense | No |
| Secondary (nonkey) | Number of records ^b or number of distinct index field values ^c | Dense or Nondense | No |

^aYes if every distinct value of the ordering field starts a new block; no otherwise.

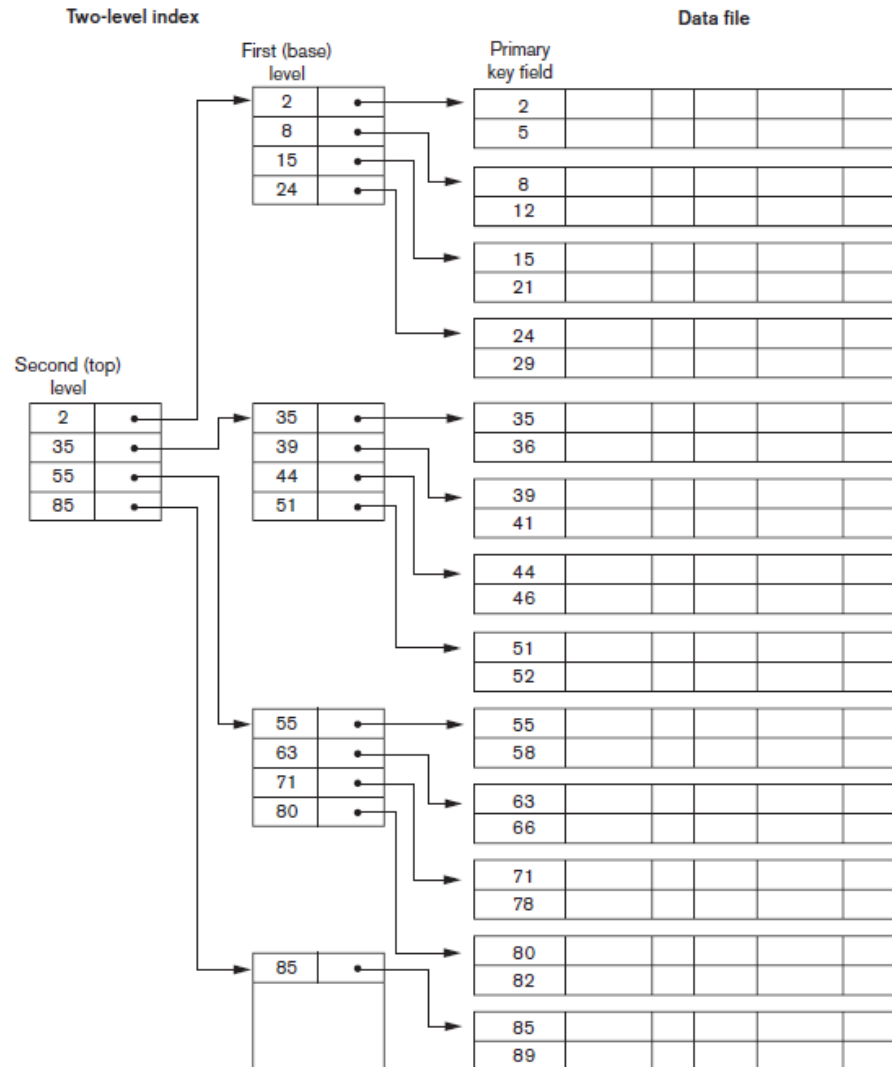
^bFor option 1.

^cFor options 2 and 3.

17.2 Multilevel Indexes

- Designed to greatly reduce remaining search space as search is conducted
- Index file
 - Considered first (or base level) of a multilevel index
- Second level
 - Primary index to the first level
- Third level
 - Primary index to the second level

Figure 17.6 A Two-Level Primary Index Resembling ISAM (Indexed Sequential Access Method) Organization




17.3 Dynamic Multilevel Indexes Using B-Trees

- Tree data structure terminology
 - Tree is formed of nodes
 - Each node (except root) has one parent and zero or more child nodes
 - Leaf node has no child nodes
 - Unbalanced if leaf nodes occur at different levels
 - Nonleaf node called internal node
 - Subtree of node consists of node and all descendant nodes

Tree Data Structure


Figure 17.7 A tree data structure that shows an unbalanced tree



Search Trees and B-Trees

- Search tree used to guide search for a record
 - Given value of one of record's fields

Figure 17.8 A node in a search tree with pointers to subtrees below it




⁸This restriction can be relaxed. If the index is on a nonkey field, duplicate search values may exist and the node structure and the navigation rules for the tree may be modified.

Search Trees and B-Trees

- Algorithms necessary for inserting and deleting search values into and from the tree

Figure 17.9 A search tree of order $p = 3$

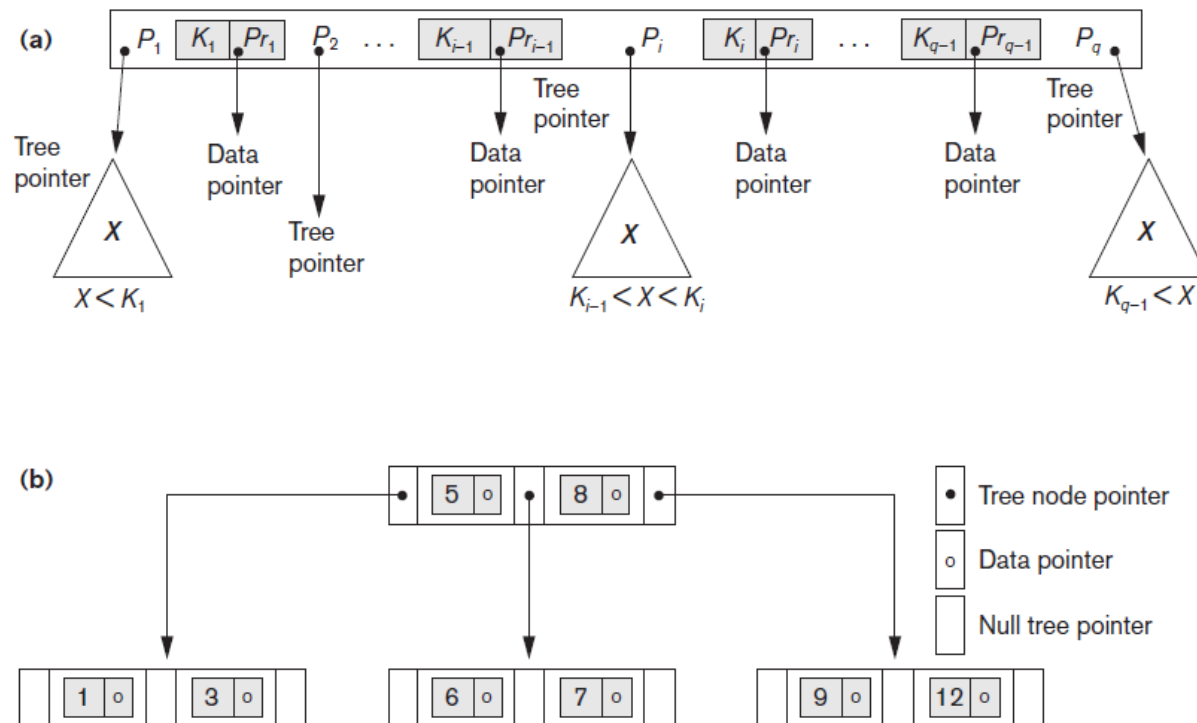


B-Trees

- Provide multi-level access structure
- Tree is always balanced
- Space wasted by deletion never becomes excessive
 - Each node is at least half-full
- Each node in a B-tree of order p can have at most $p-1$ search values

B-Tree Structures

Figure 17.10 B-tree structures (a) A node in a B-tree with $q-1$ search values (b) A B-tree of order $p=3$. The values were inserted in the order 8, 5, 1, 7, 3, 12, 9, 6



Index Creation

- General form of the command to create an index

```
CREATE [ UNIQUE ] INDEX <index name>  
ON <table name> ( <column name> [ <order> ] { , <column name> [ <order> ] } )  
[ CLUSTER ] ;
```

- Unique and cluster keywords optional
- Order can be ASC or DESC

Indexing of Strings

- Strings can be variable length
- Strings may be too long, limiting the fan-out
- Prefix compression
 - Stores only the prefix of the search key adequate to distinguish the keys that are being separated and directed to the subtree

Physical Database Design Decisions

- Design decisions about indexing
 - Whether to index an attribute
 - Attribute is a key or used by a query
 - What attribute(s) to index on
 - Single or multiple
 - Whether to set up a clustered index
 - One per table

17.8 Summary

- Indexes are access structures that improve efficiency of record retrieval from a data file
- Ordered single-level index types
 - Primary, clustering, and secondary
- Multilevel indexes can be implemented as B-trees and B+-trees
 - Dynamic structures